



Assignment #1

The Weight of Efficiency

MPG–Weight trade-off with Horsepower, Cylinders, Origin, Year, and Model encoded

Friday 31st October, 2025 15:30

Student:

Maarten Stork - 15761770

Course:

Scientific Visualization and Virtual Reality

Lecturer:

Dr. Rob Belleman

Course code:

5284SVVR6Y

Abstract

This report presents a single visualization that explains fuel efficiency (MPG) across a `cars` dataset while encoding seven different attributes. Following principles from Bertin and Mackinlay, MPG vs. weight is mapped to position×position (highest accuracy), horsepower to point size, origin to shape, cylinders to value, and year to a lightness ramp on the marker edge; model names are kept out to avoid clutter, but are present upon hover. The plot shows a strong inverse MPG-weight relationship, modulated by horsepower, with European/Japanese cars clustering at lower weight and higher MPG, and a temporal drift (1970–1982) toward better efficiency. This design maximizes readability by reserving spatial position for the main message and using separable, type-appropriate channels for the rest, ordered by level of importance.

1 Introduction

Fuel efficiency is shaped by a web of interacting factors: amongst which a car's mass and engine power, the engineering choices behind cylinder count, the era in which it was built, and regional design philosophies. Understanding how such variables jointly affect miles per gallon (MPG) informs both history and today's efficiency trade-offs - a key issue given cars' role in energy waste and greenhouse pollution [1].

This report uses a single, information-rich visualization to study all cars in a dataset (with variables *model*, *MPG*, *cylinders*, *horsepower*, *weight*, *year*, and *origin*) as supplied by the University of Amsterdam¹. The narrative is focused on fuel efficiency versus vehicle weight as mass is the main driver of energy use in conventional cars (heavier cars need more energy to move, directly lowering MPG), a relationship the U.S. EPA directly identifies as a primary determinant of fuel consumption and tailpipe CO₂ [2]. Around that core, we also examine the remaining variables in the dataset to see how they interact with weight and fuel efficiency, and whether any secondary patterns emerge that would not

¹Dataset is available here: <https://github.com/MaartenStork/SVVR/blob/main/Assignment1/cars.csv>

be obvious from the raw data alone. This method of visualization follows Bertin's semiology (visual variables and ordering) [3] and Mackinlay's expressiveness/effectiveness principles [5].

Our setup gives us the following (natural) research question:

What is the relationship between vehicle weight and fuel efficiency (MPG), and how is this relationship modulated by horsepower and cylinder count, varying across region, model year ('70-'82), and individual model?

The results will hopefully sharpen insight into where efficiency is actually won or lost, helping allocate attention (and improvements) where they matter most.

2 Methods

Data and Type Classification We analyze the provided `cars.csv` with seven attributes per car: *model*, *MPG*, *cylinders*, *horsepower*, *weight* (spelled *weighth* in the file), *year*, and *origin*. Column names are standardized (*weighth*→*weight*); numeric fields are coerced to numbers (`errors="coerce"`); and rows missing any of the variables are removed. When *year* appears as two digits (70–82), we define *year.full* = 1900 + *year* for ordered mappings. We refer readers to the Appendix code for the exact implementation details. Additionally, see Table 1 for an overview and explanation of the columns and their types in the `cars` dataset.

Table 1: *Data type classification and justification for the seven columns.*

Column	Type	Justification
<i>model</i>	Nominal	Identifier; no inherent order.
<i>MPG</i>	Quant. (ratio)	Positive, continuous; differences/ratios meaningful.
<i>cylinders</i>	Ordinal	Integer count; naturally ordered classes.
<i>horsepower</i>	Quant. (ratio)	Positive, continuous; differences/ratios meaningful.
<i>weight</i>	Quant. (ratio)	Mass proxy; differences/ratios meaningful.
<i>year</i>	Ordinal	Differences meaningful; used as ordered index.
<i>origin</i>	Nominal	Region label (US/Europe/Japan).

To align the visualization with perceptual accuracy and the assignment's requirement to encode all variables, we adopt the following importance order: (i) MPG (outcome to explain), (ii) weight (dominant physical driver), (iii) horsepower (performance factor that shifts MPG at fixed weight), (iv) cylinders (engine class proxy), (v) year (temporal drift, 1970–1982), (vi) origin (regional grouping), (vii) model (identifier). This approach lets readers first assess the core quantitative trade-off (MPG vs. weight) using the most accurate channel (position), then layer on separable channels to explain structure around that trend - ending with the least globally informative but lookup-useful detail, the model identifier for point-specific inspection. More details on this are given following the next paragraph.

Design principle and mappings Following Bertin's variables and Mackinlay's effectiveness ordering, we (i) reserve spatial position for the most important quantitative relation and (ii) assign type-appropriate, separable channels to the remaining variables so they add information without competing with the axes.

For the quantitative data, we use channels that support precise magnitude and comparison [5]. Thus, we map *MPG* to the y-position and *weight* to the x-position, so the primary quantitative trade-off is read with the most accurate channel (position×position). Axes use simple linear scales so positions read directly. *Horsepower* is then encoded by area (size): it's the strongest remaining quantitative cue, uses a monotone (linear) stretch between a small minimum (to keep low values visible) and a capped maximum (to limit overlap), and stays secondary to position while remaining

separable from other channels. We intentionally avoided a 3D plot (even with three quantitative variables lending itself to position \times position \times position) because in a static medium we believed it adds some clutter and perspective distortion, while making it more difficult to read and compare.

Then, for the nominal types, we use channels that emphasize clear grouping and comparison across categories [5]. *Origin* is encoded with three easily recognizable shapes that align with common cultural cues: Europe as a square (a simple, neutral geometric block often used for regions/blocks in maps and charts), the USA as a star (referencing the stars on the U.S. flag), and Japan as a circle (reflecting the red sun disc of the Hinomaru). Using this highly discriminable symbol set makes groups preattentively separable without implying any numeric order, keeps the legend compact, and avoids interference with other channels. Secondly, *model* is kept as on-demand tooltip/metadata for point-level lookup; the popup shows only the model name so identification stays separate from the other visuals (that are seen at a glance) and the plot remains uncluttered.

Lastly, for the ordinal datatypes, we use channels that make rank/order obvious without overclaiming precise magnitude [5]. *Year* is normalized across its observed range and encoded as a value for edge lightness (in grayscale, where earlier is darker, later lighter) so temporal order is conveyed subtly. *Cylinders* are encoded with a sequential blue-green value palette, where lighter, desaturated tints represent fewer cylinders and darker, more saturated tones represent more cylinders (3 cyl = light green, . . . , 8 cyl = deep blue-green). A sequential scale conveys order without implying precise magnitude, making it well-suited to ordinal data. We sort the legend in the order of color-values to reinforce rank, and we keep the mapping perceptually monotonic so equal steps in cylinder count produce consistent visual steps. As evident, we chose to order value scales for both *year* and *cylinders*, and we explicitly avoid using color for nominal fields. This follows accessibility literature: sequential palettes with a clear lightness change are easier to read (including for many color-blind readers [7]) and are a good fit for ordered data, whereas mixing multiple unrelated color schemes reduces distinguishability [5, 6].

Referring back to the importance order in the paragraph on **Data and Type Classification**, we made sure each variable's priority aligns with an appropriate visual channel, not just in the context of what kind of data type, but also in the context of the importance laid out in that subsection. For example, for our research question, we judged *horsepower* to be more important than *cylinders*: we center the analysis on MPG versus weight and consider horsepower a more direct design lever than cylinder count. Following Mackinlay's principle that more important features should receive more effective channels [5], we then map *horsepower* to mark size (area), making high-power cars immediately noticeable from afar, and map *cylinders* to a blue-green sequential value. The result is a clear size for horsepower, with cylinder count conveyed by a subtler, more gradual, ordered shift in value. This is one example; the same priority-driven mapping, alongside correct channel–data-type pairing, was the base for every channel-assignment in the figure.

3 Results

Figure 1 presents the complete cars dataset in a single scatter plot. Legends for value, shape, and size are shown on the right; a grayscale colorbar below the plot indicates the mapping from year to edge lightness. Axes titles and a light grid for better positional reading.

Figure 2 shows a magnified excerpt of the same plot illustrating the on-hover tooltip, which reveals the corresponding *model* name.

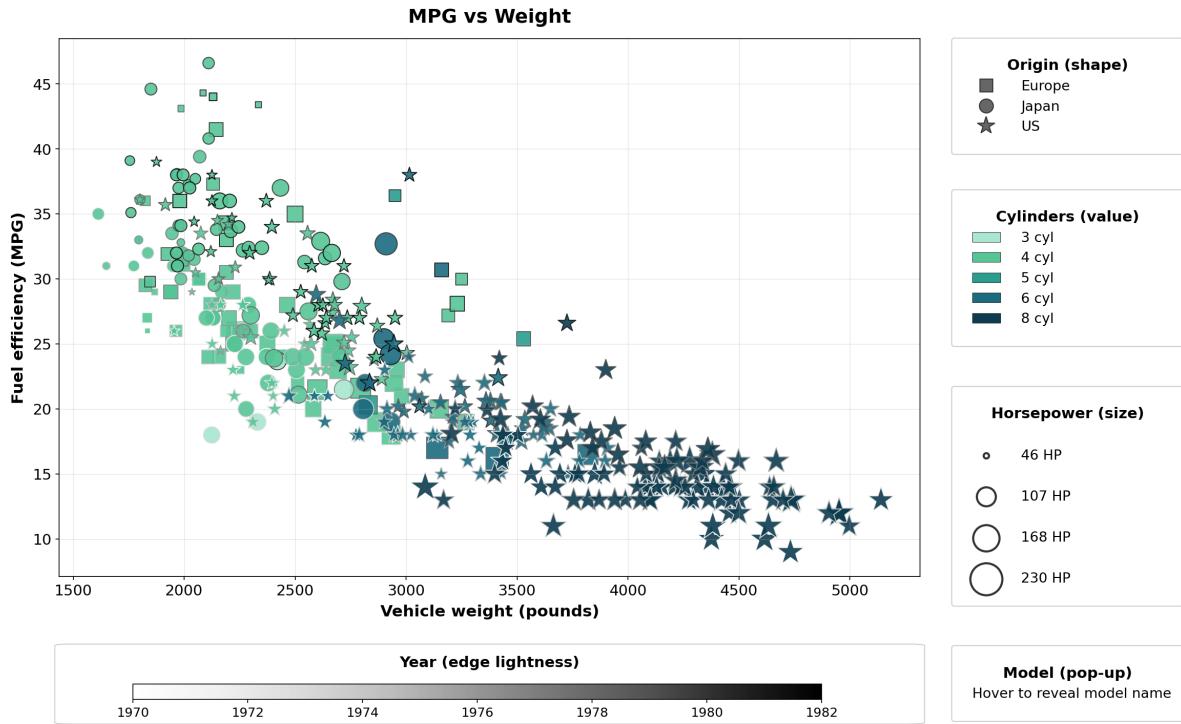


Figure 1: Single-figure encoding of all seven attributes: weight (x), MPG (y), horsepower (size), origin (shape), cylinders (value), and year (edge lightness); model names are available on hover in the SVG export.



Figure 2: Detail view showing the on-hover model label in the SVG version of the figure (example: “Chevrolet Citation”).

4 Discussion

The visualization answers the primary question clearly: there is a strong inverse relationship between vehicle weight and fuel efficiency (MPG). Cars cluster along a downward trend (i.e. heavier vehicles achieve fewer miles per gallon) making mass a very dominant determinant of efficiency in this dataset. This core pattern is legible at a glance because it is encoded with position \times position, the most accurate channel.

Secondary patterns help explain other notabilities around that trend. Larger markers (more horsepower) tend to sit lower on the MPG axis with higher weights, consistent with what we expect of performance trading off against efficiency. Cylinder count also aligns with this: higher-cylinder engines appear mostly in the heavier, lower-MPG region, while low-cylinder engines are concentrated among lighter, more efficient cars (virtually none among the heaviest vehicles). Regional groupings show many U.S. models in the heavier/less efficient space, while European and Japanese models

more often occupy the lighter/higher-MPG zone. The grayscale edge encodes time and indicates a temporal drift from 1970 to 1982 toward lighter, more efficient cars with fewer cylinders and lower horsepower. This is also consistent with known information on the period's regulatory and market pressures [4]. Although model names add little to answering the research question, the option to show them does support identifying specific outliers and inviting comparisons.

References

- [1] Jillian Anable and Christian Brand. Energy, pollution and climate change. In *Transport Matters*, pages 55–82. Policy Press, 2019.
- [2] Edward Becker. Automotive trends report. *Tribology & Lubrication Technology*, 80(6):96–98, 2024.
- [3] Jacques Bertin. *Semiology of Graphics: Diagrams, Networks, Maps*. University of Wisconsin Press, Madison, 1983.
- [4] Melvyn A Fuss and Leonard Waverman. Productivity growth in the automobile industry, 1970-1980: A comparison of canada, japan and the united states. Technical report, National Bureau of Economic Research, 1985.
- [5] Jock Mackinlay. Automating the design of graphical presentations of relational information. *Acm Transactions On Graphics (Tog)*, 5(2):110–141, 1986.
- [6] Kenneth Moreland. Diverging color maps for scientific visualization. In *International symposium on visual computing*, pages 92–103. Springer, 2009.
- [7] Christine Rigden. 'the eye of the beholder'-designing for colour-blind users. *British Telecommunications Engineering*, 17:291–295, 1999.

Appendix

For completeness, I also showcase the core script below. The complete repository (containing the code, `cars` dataset, and several alternative visualizations) is available on GitHub².

Listing 1: Data cleaning + visualization script

```

1  # cars_viz_vibrant.py
2
3  import math
4  import pandas as pd
5  import numpy as np
6  import matplotlib.pyplot as plt
7  from matplotlib.lines import Line2D
8  from matplotlib.patches import Patch
9  import matplotlib as mpl
10 import mplcursors
11
12 # ----- Load & clean -----
13 df = pd.read_csv("cars.csv")
14
15 # Normalize column names (handle misspelling "weigth")

```

²See the full codebase here: <https://github.com/MaartenStork/SVVR/main/Assignment1>

```
16 cols = {c.lower().strip(): c for c in df.columns}
17 # rename to consistent lower-case names
18 rename_map = {}
19 if "weigth" in cols: # original misspelling
20     rename_map[cols["weigth"]] = "weight"
21 if "mpg" in cols:
22     rename_map[cols["mpg"]] = "mpg"
23 if "cylinders" in cols:
24     rename_map[cols["cylinders"]] = "cylinders"
25 if "horsepower" in cols:
26     rename_map[cols["horsepower"]] = "horsepower"
27 if "year" in cols:
28     rename_map[cols["year"]] = "year"
29 if "model" in cols:
30     rename_map[cols["model"]] = "model"
31 if "origin" in cols:
32     rename_map[cols["origin"]] = "origin"
33
34 df = df.rename(columns=rename_map)
35
36 # Coerce numeric columns; some datasets store '?' for horsepower.
37 for col in ["mpg", "weight", "horsepower", "year", "cylinders"]:
38     if col in df.columns:
39         df[col] = pd.to_numeric(df[col], errors="coerce")
40
41 # If years are 70..82, convert to full years 1970..1982 for labeling; keep numeric continuity.
42 if df["year"].dropna().between(0, 150).all(): # heuristic for model-year encoding
43     df["year_full"] = np.where(df["year"] < 100, 1900 + df["year"], df["year"])
44 else:
45     df["year_full"] = df["year"]
46
47 # Drop rows missing essentials for the mark
48 df = df.dropna(subset=["mpg", "weight", "horsepower", "cylinders", "origin", "year_full"]).copy()
49
50 # ----- Encodings -----
51 # Position -> (weight, mpg)
52 x = df["weight"].values
53 y = df["mpg"].values
54
55 # Size (area) -> horsepower (use perceptual area scaling)
56 hp = df["horsepower"].values.astype(float)
57 hp_min, hp_max = float(np.nanmin(hp)), float(np.nanmax(hp))
58 # Map horsepower to area in [Amin, Amax]; choose radii so small values are still visible.
59 Amin, Amax = 25.0, 900.0 # scatter 's' argument is area in points^2
60 hp_norm = (hp - hp_min) / (hp_max - hp_min + 1e-9)
61 sizes = Amin + hp_norm * (Amax - Amin)
62
63 # Shape -> origin (nominal)
64 # Different shapes for each origin
65 origin_markers = {
66     'US': '*', # star
67     'Europe': 's', # square
68     'Japan': 'o' # circle
69 }
70
```

```
71 def marker_for_origin(o):
72     return origin_markers.get(str(o), "o")
73
74 origins = df["origin"].astype(str)
75 origin_levels = list(pd.Categorical(origins).categories)
76
77 # Color (facecolor) -> cylinders (ordered - sequential color scale)
78 # Map cylinders to colors - darker as more cylinders
79 unique_cyls = sorted(df["cylinders"].dropna().unique())
80 # Create a blue-green sequential palette (lighter to darker)
81 from matplotlib.colors import LinearSegmentedColormap
82 cyl_colors = {
83     3: '#a8e6cf', # Lightest green
84     4: '#56c596', # Light green
85     5: '#2a9d8f', # Medium teal
86     6: '#1b6b7f', # Dark teal
87     8: '#0d3b4f' # Darkest blue-green
88 }
89
90 def color_for_cyl(c):
91     return cyl_colors.get(int(c), '#2a9d8f')
92
93 # Year (ordered) -> lightness on the marker EDGE using a grayscale ramp
94 yr = df["year_full"].values.astype(float)
95 yr_min, yr_max = float(yr.min()), float(yr.max())
96 yr_norm = (yr - yr_min) / (yr_max - yr_min + 1e-9)
97
98 # Build a grayscale mapper for edges (0 -> dark, 1 -> light)
99 year_cmap = mpl.colormaps.get_cmap("Greys")
100 edge_colors = [year_cmap(val) for val in yr_norm]
101
102 # ----- Figure -----
103 # Create figure with more space for legends and margins
104 fig = plt.figure(figsize=(16, 10), dpi=150)
105
106 # Set default font size larger
107 plt.rcParams.update({'font.size': 12})
108
109 # Create main plot area with specific position to leave room for legends
110 ax = plt.axes([0.06, 0.15, 0.70, 0.70]) # [left, bottom, width, height]
111
112 # Store all points and their model names for hover functionality
113 all_scatter_objects = []
114 model_names = []
115
116 # Plot by origin (for shape) and cylinders (for color)
117 for orig in origin_levels:
118     subset = df[origins == orig]
119     if subset.empty:
120         continue
121
122     # Get the marker for this origin
123     marker = marker_for_origin(orig)
124
125     # Now group by cylinders to get colors
```

```
126     for cyl_val in sorted(subset["cylinders"].dropna().unique()):  
127         cyl_subset = subset[subset["cylinders"] == cyl_val]  
128         if cyl_subset.empty:  
129             continue  
130  
131         idx = cyl_subset.index  
132  
133         # Get color for this cylinder count  
134         fc = color_for_cyl(cyl_val)  
135  
136         # Get sizes for this subset  
137         ss = sizes[idx]  
138  
139         # Get edge colors for year  
140         ec = [year_cmap(v) for v in ((cyl_subset["year_full"].values.astype(float) - yr_min) / (yr_max -  
141             yr_min + 1e-9))]  
142  
143         # Get data  
144         sub_x = cyl_subset["weight"].values  
145         sub_y = cyl_subset["mpg"].values  
146         sub_models = cyl_subset["model"].astype(str).values  
147  
148         sc = ax.scatter(sub_x, sub_y, s=ss, marker=marker,  
149                         facecolor=fc, edgecolor=ec, linewidth=0.9, alpha=0.9)  
150         all_scatter_objects.append(sc)  
151         # Store model names for this scatter object  
152         model_names.append(sub_models.tolist())  
153  
154     # Add hover tooltips using mplcursors with beautiful styling  
155     def make_annotation_func(scatter_objs, model_lists):  
156         """Create annotation function that shows model name on hover"""  
157         def annotate(sel):  
158             # Find which scatter object this point belongs to  
159             point_idx = sel.index  
160             for scatter_obj, models in zip(scatter_objs, model_lists):  
161                 if sel.artist == scatter_obj:  
162                     if point_idx < len(models):  
163                         model = models[point_idx]  
164                         # Normalize model name - capitalize properly  
165                         model_clean = model.strip().title()  
166                         # Show only the model name  
167                         sel.annotation.set_text(model_clean)  
168                         sel.annotation.set_fontsize(13)  
169                         sel.annotation.set_fontweight('600') # Semi-bold  
170  
171                         # Style the box beautifully  
172                         bbox = sel.annotation.get_bbox_patch()  
173                         bbox.set_boxstyle("round,pad=0.7")  
174                         bbox.set_facecolor('#2c3e50') # Dark blue-gray  
175                         bbox.set_edgecolor('#3498db') # Bright blue border  
176                         bbox.set_linewidth(2)  
177                         bbox.set_alpha(0.95)  
178  
179                         # Style the text  
180                         sel.annotation.set_color('white')
```

```
180
181         # Style the arrow
182         sel.annotation.arrow_patch.set_color('#3498db')
183         sel.annotation.arrow_patch.set_linewidth(1.5)
184     break
185 return annotate
186
187 # Create cursor with hover functionality
188 cursor = mplcursors.cursor(all_scatter_objects, hover=True)
189 cursor.connect("add", make_annotation_func(all_scatter_objects, model_names))
190
191 # Axes labels & title
192 ax.set_xlabel("Vehicle weight (pounds)", labelpad=6, fontsize=15, fontweight='bold')
193 ax.set_ylabel("Fuel efficiency (MPG)", labelpad=6, fontsize=15, fontweight='bold')
194
195 # Make tick labels bigger
196 ax.tick_params(axis='both', which='major', labelsize=14)
197 ax.set_title("MPG vs Weight",
198             fontsize=18, pad=15, fontweight='bold')
199
200 # Make grid subtle to support reading by position
201 ax.grid(True, linestyle="--", linewidth=0.5, alpha=0.4)
202
203 # ----- Legends (Unified Professional Layout) -----
204
205 # Create unified legend area on the right
206 legend_x = 0.78 # X position for all legends
207 legend_width = 0.20 # Width for legends
208 legend_height = 0.16 # Height for each legend box
209 legend_gap = 0.09 # Consistent gap between all legends
210 legend_gap_large = 0.178 # Larger gap before Horsepower
211
212 # Style settings for consistency
213 legend_style = {
214     'frameon': True,
215     'framealpha': 1.0,
216     'edgecolor': '#d0d0d0',
217     'fancybox': True,
218     'fontsize': 13,
219     'title_fontsize': 14,
220     'borderpad': 1.5,
221     'columnspacing': 1.0,
222     'handletextpad': 1.5
223 }
224
225 # Calculate positions with proper spacing
226 legend_top = 0.69 # Top position for first legend
227
228 # 1. Origin legend - Top (now shapes)
229 # Adjust marker sizes so they appear visually similar
230 marker_sizes = {
231     'Europe': 12, # square
232     'Japan': 13, # circle (slightly bigger)
233     'US': 17 # star (needs to be bigger to appear same size)
234 }
```

```
235 origin_handles = [Line2D([0], [0], marker=origin_markers[o], color="none",
236                         markerfacecolor="#666666", markeredgecolor="#333333",
237                         markersize=marker_sizes.get(o, 12), linewidth=0)
238                         for o in origin_levels]
239 origin_y = legend_top
240 origin_ax = fig.add_axes([legend_x, origin_y, legend_width, legend_height])
241 origin_ax.axis('off')
242 leg1 = origin_ax.legend(origin_handles, origin_levels, title="Origin (shape)",
243                         loc='lower left', mode='expand', ncol=1,
244                         bbox_to_anchor=(0, 0, 1, 1),
245                         **legend_style)
246 leg1.get_title().set_fontweight('bold')
247
248 # 2. Cylinders legend - Middle (now colors)
249 unique_cyl = sorted(df["cylinders"].dropna().astype(int).unique())
250 cyl_handles = [Patch(facecolor=cyl_colors[c], edgecolor="#333333", linewidth=0.5)
251                 for c in unique_cyl]
252 cylinders_y = origin_y - legend_height - legend_gap
253 cylinders_ax = fig.add_axes([legend_x, cylinders_y, legend_width, legend_height])
254 cylinders_ax.axis('off')
255 leg2 = cylinders_ax.legend(cyl_handles, [f"{c} cyl" for c in unique_cyl],
256                             title="Cylinders (value)", loc='lower left', mode='expand', ncol=1,
257                             bbox_to_anchor=(0, 0, 1, 1),
258                             **legend_style)
259 leg2.get_title().set_fontweight('bold')
260
261 # 3. Horsepower legend - Bottom
262 hp_ticks = np.linspace(hp_min, hp_max, 4)
263 hp_sizes = Amin + (hp_ticks - hp_min) / (hp_max - hp_min + 1e-9) * (Amax - Amin)
264 size_handles = [plt.scatter([], [], s=s, facecolor="white", edgecolor="#333333",
265                             linewidth=2.0)
266                             for s in hp_sizes]
267 hp_y = cylinders_y - legend_height - legend_gap_large
268 hp_ax = fig.add_axes([legend_x, hp_y, legend_width, legend_height])
269 hp_ax.axis('off')
270 leg3 = hp_ax.legend(size_handles, [f"{int(v)} HP" for v in hp_ticks],
271                      title="Horsepower (size)", loc='lower left', mode='expand', ncol=1,
272                      bbox_to_anchor=(0, 0, 1, 1),
273                      scatterpoints=1, labelspacing=2.0,
274                      **legend_style)
275 leg3.get_title().set_fontweight('bold')
276
277 # 4. Model legend (hover info) - Below Horsepower
278 from matplotlib.patches import Rectangle
279 model_gap = -0.01 # Smaller gap for Model legend
280 model_y = hp_y - legend_height - model_gap
281 model_ax = fig.add_axes([legend_x, model_y, legend_width, legend_height])
282 model_ax.axis('off')
283
284 # Create invisible handles for the text
285 invisible_handle = Rectangle((0, 0), 0, 0, alpha=0)
286 legend_style_no_handle = {k: v for k, v in legend_style.items()
287                           if k not in ['handletextpad', 'handlelength']}
288 leg4 = model_ax.legend([invisible_handle],
289                       ['Hover to reveal model name'],
```

```
290         title="Model (pop-up)",
291         loc='lower left', mode='expand', ncol=1,
292         bbox_to_anchor=(0, 0, 1, 1),
293         handlelength=0,
294         handletextpad=0,
295         **legend_style_no_handle)
296 leg4.get_title().set_fontweight('bold')
297
298 # Year colorbar - with proper box styling
299 norm = mpl.colors.Normalize(vmin=yr_min, vmax=yr_max)
300 sm = mpl.cm.ScalarMappable(cmap=year_cmap, norm=norm)
301 sm.set_array([])
302
303 # First create the box background - slightly taller and moved down more
304 box_ax = fig.add_axes([0.07, -0.04, 0.68, 0.10])
305 box_ax.set_xlim(0, 1)
306 box_ax.set_ylim(0, 1)
307 box_ax.axis('off')
308
309 # Draw a rectangle with rounded corners matching the legend style
310 from matplotlib.patches import FancyBboxPatch
311 rect = FancyBboxPatch(
312     (0, 0), 1, 1,
313     boxstyle="round,pad=0.02",
314     facecolor='white',
315     edgecolor="#d0d0d0",
316     linewidth=1,
317     transform=box_ax.transAxes,
318     clip_on=False,
319     zorder=1 # Behind colorbar
320 )
321 box_ax.add_patch(rect)
322
323 # Add the title
324 box_ax.text(0.5, 0.87, 'Year (edge lightness)', fontsize=14, fontweight='bold',
325             ha='center', va='top', transform=box_ax.transAxes, zorder=3)
326
327 # Now create the colorbar ON TOP of the box - adjusted for new box position
328 cbar_ax = fig.add_axes([0.12, -0.01, 0.56, 0.02])
329 cbar = plt.colorbar(sm, cax=cbar_ax, orientation='horizontal')
330 cbar.ax.tick_params(labelsize=12)
331 cbar.ax.set_zorder(2) # On top of box
332
333 # ----- Nice ticks -----
334 # If year_full was 1970..1982, ensure the colorbar shows integers
335 if yr_max - yr_min <= 20:
336     cbar.set_ticks(np.arange(math.floor(yr_min), math.ceil(yr_max)+1, 2))
337
338 # No need for tight_layout since we're manually positioning everything
339
340 # Save both PNG and SVG with margins to ensure nothing is cut off
341 plt.savefig("cars_viz_vibrant_final.png", dpi=150, bbox_inches='tight', pad_inches=0.5)
342 plt.savefig("cars_viz_vibrant_final.svg", bbox_inches='tight', pad_inches=0.5)
343
344 print(" Saved: cars_viz_vibrant_final.png")
```

```
345 print(" Saved: cars_viz_vibrant_final.svg")
346
347 plt.show()
```