



Assignment #2

The Square That Sets the Pace

Effect of hot-square size on Jacobi convergence in a 2-D Laplace problem

Friday 7th November, 2025 16:55

Student:

Maarten Stork - 15761770

Course:

Scientific Visualization and Virtual Reality

Lecturer:

Dr. Rob Belleman

Course code:

5284SVVR6Y

Abstract

This report studies how the size of a central hot square affects the convergence of a Jacobi solver for a 2-D steady-state heat problem. Using the by-UvA-provided setup, we solve Laplace's equation on a 9×9 plate with Dirichlet boundaries (cool bottom, warm top, linear side ramps) and a fixed-temperature heater at the center. The solver is a standard Jacobi method with a five-point stencil and a max-norm δ stopping rule. We vary only the heater area (as a fraction of the domain) and compare both the iterations needed to reach the tolerance and the full δ -vs-iteration history. Results are shown in two ways: (i) ParaView outputs (VTI frames with a PVD time series) documented via a Haber&McNabb pipeline, and (ii) Matplotlib figures/animations for a lightweight web companion. Across the tested sizes, larger interior Dirichlet regions converge in fewer iterations.

1 Introduction

Laplace's equation sits at the heart of many steady-state phenomena, among which heat flow, electrostatics, and potential fluid models [2, 4]. Its solutions are entirely shaped by what we fix on the boundary (or inside the domain), so simple geometric changes can alter both the final field and how quickly an iterative solver gets there. Jacobi's method, while basic, is ideal for seeing this link between problem geometry and convergence behaviour, and scientific visualization makes that link easily detectable by showing the field evolve under consistent scales and annotations [1, 5, 6].

We adopt the UvA "cartoon" setup ¹: a square plate with a cool bath at one edge, a warm edge opposite, linear side ramps, and a fixed-temperature hot square embedded at the center (think of a small heater inside the material). Starting from a plain initial guess, the system relaxes to steady state. In this study we vary only one element (the area

¹Access to the full cartoon through this Github repository: <https://github.com/MaartenStork/SVVR/blob/main/GeneralPlots/CartoonUVA.pdf>

of the embedded heater) and examine how that modelling choice affects solver effort, giving us the following research question:

How does the size of the central hot square affect the number of iterations a Jacobi solver needs to reach a given tolerance?

The aim is to provide a clear, visual account of convergence alongside a simple quantitative comparison across heater sizes, keeping all other modelling and visualization choices fixed.

2 Methods

Problem setup We model steady heat conduction on a square plate: the bottom edge is held cool ($T=32^\circ$), the top edge warm ($T=100^\circ$), and the left/right edges ramp linearly between them ($T=32$ to 100°); a hot square ($T=212^\circ$) embedded at the center is fixed at a higher temperature (an interior Dirichlet region). In steady state the temperature field $T(x, y)$ satisfies Laplace's equation $\nabla^2 T = 0$ on the plate outside the hot square, with Dirichlet conditions on all outer edges and on the embedded square. This follows the computational science pipeline from physical model → mathematical model → numerical model.

We discretize the domain on a uniform lattice (i, j) and solve the 2-D Laplace problem with the standard five-point stencil, updating only non-Dirichlet nodes. The Jacobi update is

$$T_{i,j}^{(n+1)} = \frac{1}{4} \left(T_{i-1,j}^{(n)} + T_{i+1,j}^{(n)} + T_{i,j-1}^{(n)} + T_{i,j+1}^{(n)} \right),$$

applied wherever the node is not fixed; boundary nodes and the hot-square mask are copied through each sweep. Convergence is monitored with the max-norm

$$\delta^{(n+1)} = \max_{i,j} |T_{i,j}^{(n+1)} - T_{i,j}^{(n)}|,$$

and the iteration stops when $\delta \leq \varepsilon$ (tolerance) or a safety cap of 20,000 iterations is reached.

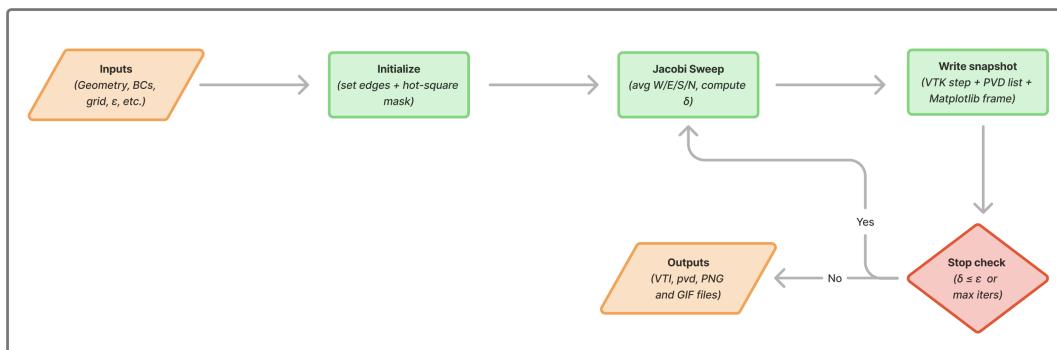


Figure 1: Flowchart (or pipeline) of the data gathering process.

This implementation (almost completely) mirrors the Jacobi pseudocode provided by the University of Amsterdam² and is depicted in Figure 2. The steps are: (1) initialize the lattice and set Dirichlet values on the outer edges as well as on the central hot-square mask; (2) iterate Jacobi sweeps using west/east/north/south neighbours while tracking δ each sweep; (3) output a snapshot at the chosen cadence (for ParaView: .vti plus a .pvd time collection); and (4)

²Readable in the GitHub at <https://github.com/MaartenStork/SVVR/blob/main/GeneralPlots/slides.pdf>.

stop when δ falls below the tolerance. As stated, we follow the slides' Jacobi loop, but instead use Dirichlet edges (no wrapping) to match the fixed-temperature boundary cartoon better; inside the loop we "clamp" neighbor indices so they never go outside the grid. This doesn't change the result, because all boundary cells already have fixed temperatures (Dirichlet). So if a neighbor would fall off the edge, we just read the boundary value that's fixed anyway. The core code of this implementation is also shown in the Appendix.

Visualization pipeline We provide two complementary visualizations, with different levels of ease:

1. For the web companion³, we render Matplotlib frames and combine these in a GIF with a fixed colormap range, axes/ticks, legend, and iteration labels; this lightweight view is meant for quick exploration and for embedding on the page, for the reviewer of this paper to run the simulations in a vacuum.
2. For the full pipeline in ParaView, the solver writes VTK-XML ImageData files (.vti) together with a .pvda collection so the sequence loads as a single time-aware source. In ParaView, the pipeline follows the Haber&McNabb model [3], also depicted in Figure 2: (1) the data source obtained through the code (flowchart described in Figure 2), (2) the mapping, as well as the additions of the contour and the annotating of time, and finally (3) rendering it all within Paraview.

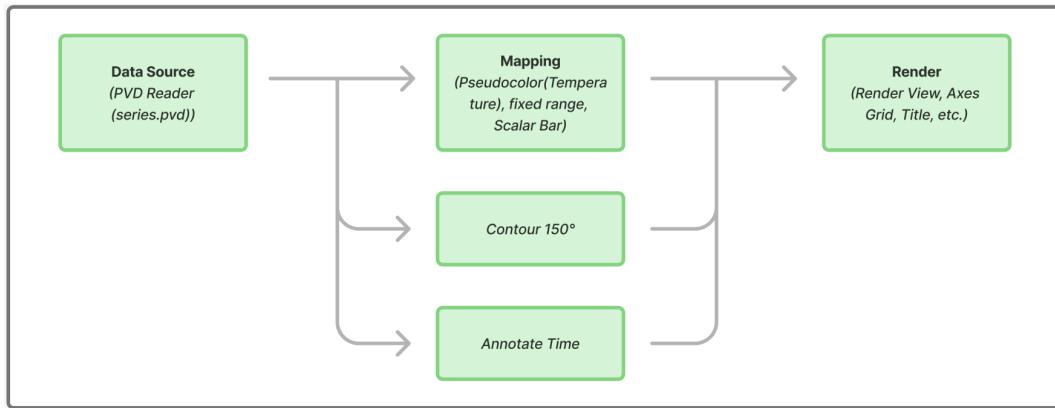


Figure 2: Pipeline of the data visualization process.

This reproduces the core visualization pipeline while offering both an accessible web version and a richer ParaView analysis. [3].

Visualization Choices Here we provide a few explanations for the visualization choices that went in to designing the plots presented in the Results section.

1. Firstly, the actual visualization. Both Matplotlib and ParaView use a physical $x, y \in [0, 9]$ domain to match the $9m \times 9m$ plate, with axes, ticks, and labels shown to give a sense of scale. We keep a single, fixed blue to red colormap range across both visualizations; a scalar bar (vertical, on the right) reinforces the mapping. In Matplotlib we use `turbo`; its hues differ slightly from the slide/ParaView palette (with a brighter yellow and green in there), and given the deadline on this particular assignment, we did not define a custom LUT. Time/iteration appears at the top of each frame for consistency. In ParaView we display it via `Annotate Time`; in Matplotlib it is rendered in the title - numerically they are the same (e.g., Matplotlib "iteration 220" corresponds to ParaView

³See the website here, feel free to play around with it: <https://jacobi-simulation-backend.onrender.com/>.

“time = 220”). For the Paraview figures we overlay a single *Contour* at a representative isotherm to indicate warm regions.

- Secondly, the plots meant for more insight into convergence behavior. For plotted convergence speed, we decide to use blue points (like our main theme for the snapshots of the visualization), with no lines inbetween, as we do not want to imply some form of linearity. For the delta-versus-iteration plot, we plot them in different values, shifting the value from purple to yellow, following Mackinlay’s effectiveness ordering for ordered data [5].

Answering the research question We run a controlled sweep over the hot-square size $f \in (0, 1]$ (fraction of plate side) in steps of $\Delta f = 0.05$, keeping all other settings fixed (geometry, grid, boundary conditions). For each f , we run the Jacobi solver to a max-norm tolerance of $\varepsilon = 10^{-4}$ and record (i) the iterations to convergence $N(f)$ and (ii) the full δ -versus-iteration history. Because the system is deterministic, a single run per f suffices. We then plot $N(f)$ to read how convergence time varies with heater size, and use the δ curves (together with consistent visualization settings) to confirm the settling behaviour across conditions.

3 Results

Visualizations Figure 3 shows the ParaView render for some of the first, middle and last frames. The contour described in the methods is visible. These are just single snapshots - a GIF of this visualization is available in the resources added to this report⁴ for the inspection of the entire animation.

Figure 4 shows the Matplotlib rendering, which one can also play around with in our online environment referenced before.

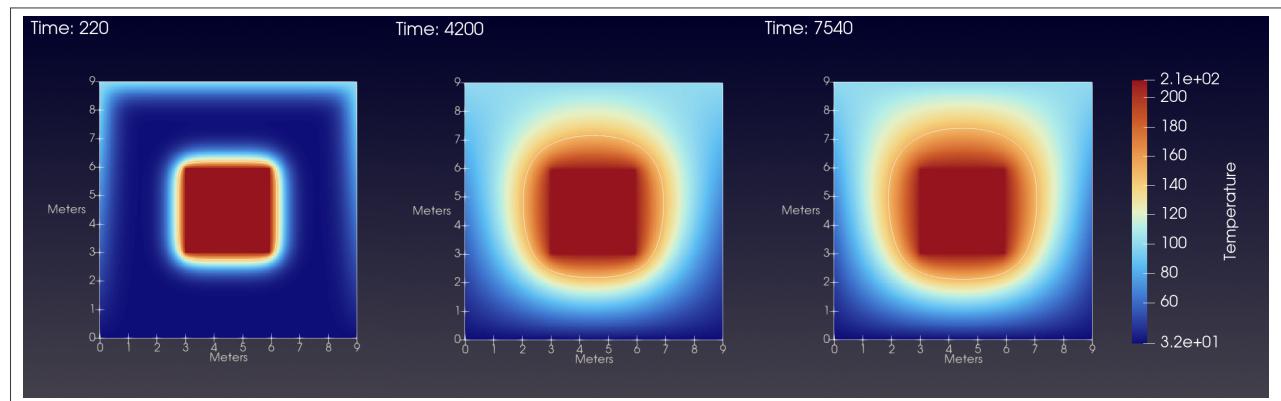


Figure 3: Three snapshots of the ParaView visualization: early, intermediate, and late stages of the temperature evolution.

Hot-square tests Figure 5 showcases several hot-square sizes to illustrate how the central heater actually changes the entire field. The side-by-side views are meant to give an intuitive feel for the experiment before we turn to the convergence results.

Figures 6 and 7 summarize solver behaviour. Figure 6 shows the speed (iterations required to reach the tolerance) so differences are visible at a glance. Figure 7 shows the history (the max-norm δ versus iteration), for each hot-square size.

⁴See the GIF in the Github here: <https://github.com/MaartenStork/SVVR/blob/main/GeneralPlots/ParaViewGIF.gif>, as well as the AVI file here: <https://github.com/MaartenStork/SVVR/blob/main/GeneralPlots/ParaViewVideo.avi>.

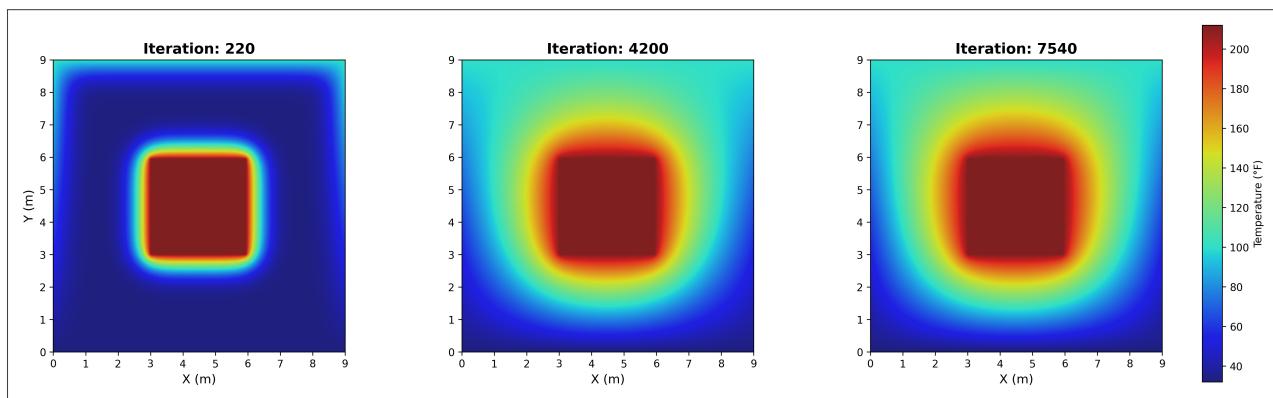


Figure 4: Three snapshots of the Matplotlib visualization: early, intermediate, and late stages of the temperature evolution. A simpler version of the ParaView visualization, meant for easy experimentation online.

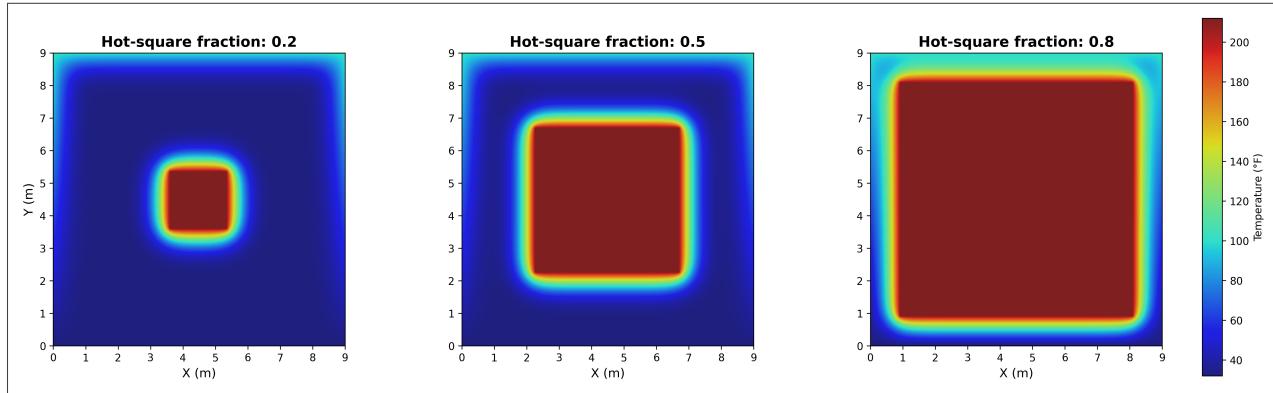


Figure 5: Three snapshots at iteration 200 with different sizes depicted in the Matplotlib visualization. Left to right; sizes 0.2, 0.5 and 0.8 respectively.

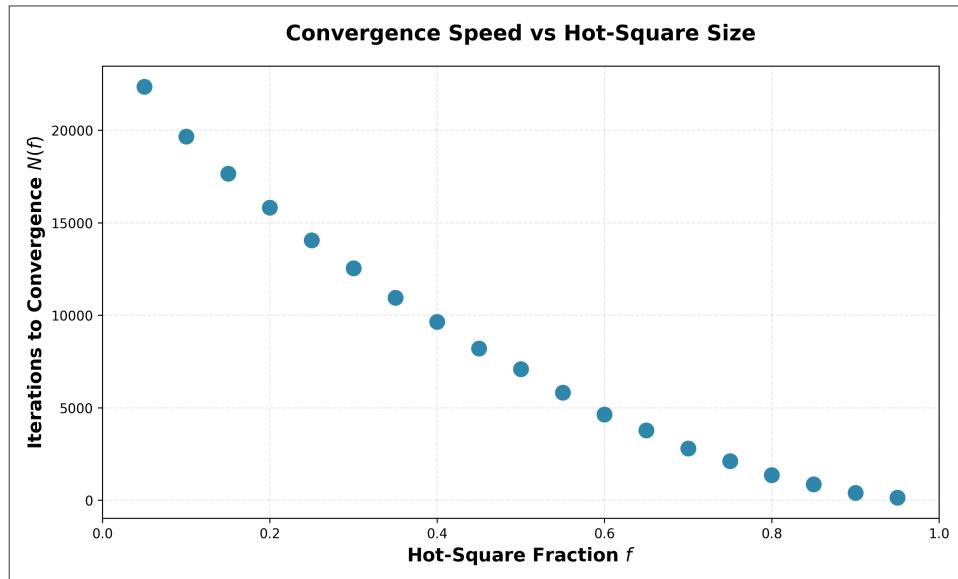


Figure 6: The course of the convergence speed against the hot-square size.

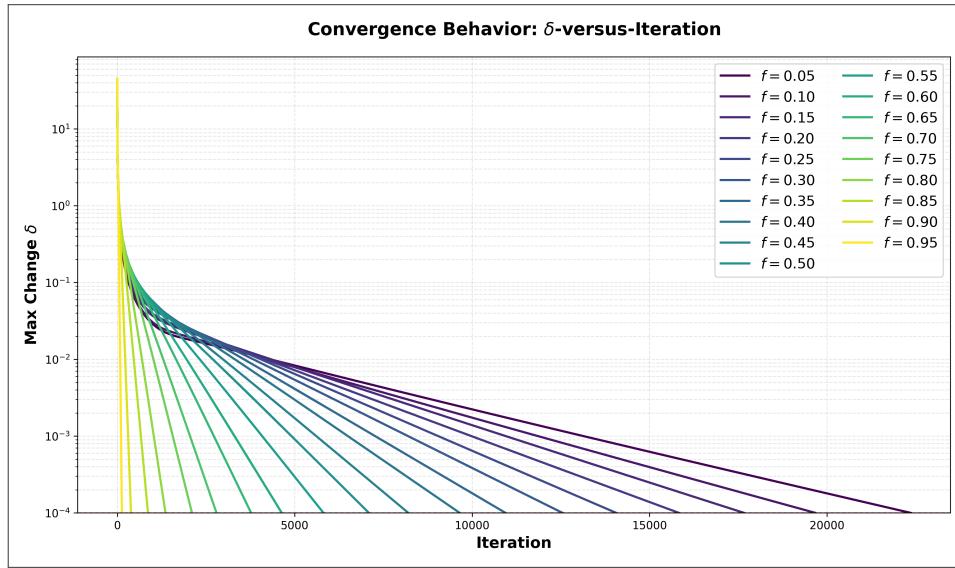


Figure 7: The δ -versus-iteration behavior for different hot-square sizes, with ' f ' meaning 'hot-square fraction'.

4 Discussion

What we see in the visualizations. With a fixed color range and a single $T=150^\circ$ contour, the early frames show a small hot patch that grows and smooths as the solver runs. Even though the heater is a square, the isotherms quickly look rounded: outside the fixed patch, the solution to Laplace's equation smooths corners and the heat spreads evenly in all directions, so level sets bend into curves rather than keeping a square outline. In the late frames the field is wide and almost concentric around the heater. Larger heaters take over more of the plate earlier, leaving a smaller region that still needs to relax, and the steady pattern appears sooner.

What we see in the graphs. The numbers match what we see in pictures and videos. The iterations-to-tolerance curve $N(f)$ drops as the heater fraction f grows: bigger heated regions need fewer sweeps. The δ -versus-iteration curves all fall quickly at first and then more slowly; on a semilog plot they look roughly linear in the tail. Smaller heaters shift these curves to the right (longer tails), while larger heaters pull them left (shorter tails). Intuitively, increasing f both reduces how many unknown nodes are left and shortens the distance over which corrections must diffuse, so the solver finishes in fewer iterations.

5 Conclusion

These results have given us an answer to our main research question:

How does the size of the central hot square affect the number of iterations a Jacobi solver needs to reach a given tolerance?

Holding geometry, grid, and tolerance fixed, larger interior Dirichlet regions (= bigger hot squares) lead to faster convergence, i.e., fewer Jacobi iterations to reach the max-norm threshold. This trend is consistent across the sweep of heater sizes and is supported both by the iteration counts and by the shape of the δ -history curves.

Looking ahead, this same setup can now serve as a testbed: we can compare different solvers, vary the heater's shape and position, or mix boundary types. With the base in place, exploring these options becomes straightforward.

References

- [1] Jacques Bertin. *Semiology of Graphics: Diagrams, Networks, Maps*. University of Wisconsin Press, Madison, 1983.
- [2] Jean Baptiste Joseph Fourier. *Théorie analytique de la chaleur*. Gauthier-Villars et fils, 1888.
- [3] Robert B Haber and David A McNabb. Visualization idioms: A conceptual model for scientific visualization systems. *Visualization in scientific computing*, 74:93, 1990.
- [4] John David Jackson and Ronald F Fox. Classical electrodynamics, 1999.
- [5] Jock Mackinlay. Automating the design of graphical presentations of relational information. *Acm Transactions On Graphics (Tog)*, 5(2):110–141, 1986.
- [6] Yousef Saad. Iterative methods for linear systems of equations: A brief historical journey. *arXiv preprint arXiv:1908.01083*, 2019.

Appendix

For completeness, I also showcase the core script below. The complete repository (containing the experimental setup, the code serving the webapp, the data, the cartoon, and more) is available on GitHub⁵.

Listing 1: Data cleaning + visualization script

```

1  from __future__ import annotations
2  import argparse
3  import math
4  import os
5  from pathlib import Path
6  from typing import Tuple, List
7  import matplotlib.pyplot as plt
8  import matplotlib.animation as animation
9  from matplotlib.colors import LinearSegmentedColormap
10 import numpy as np
11
12 # -----
13 # Utilities
14 # -----
15
16 def linspace(n: int, start: float, end: float) -> List[float]:
17     if n == 1:
18         return [start]
19     step = (end - start) / (n - 1)
20     return [start + i * step for i in range(n)]
21

```

⁵See the full codebase here: <https://github.com/MaartenStork/SVVR>

```
22 def write_vti_image_data(
23     filename: str,
24     grid: "list[list[float]]",
25     dx: float,
26     dy: float,
27     origin: Tuple[float, float, float] = (0.0, 0.0, 0.0),
28     name: str = "Temperature",
29 ):
30     """
31     Write 2D scalar field as VTK-XML ImageData (.vti), ASCII.
32     """
33     ny = len(grid)
34     nx = len(grid[0])
35     nz = 1
36     # VTK-XML expects x fastest, then y, then z (same order you already write)
37     with open(filename, "w") as f:
38         f.write('<?xml version="1.0"?>\n')
39         f.write('<VTKFile type="ImageData" version="0.1" byte_order="LittleEndian">\n')
40         f.write(f' <ImageData WholeExtent="0 {nx-1} 0 {ny-1} 0 {nz-1}" '
41                 f'Origin="{origin[0]} {origin[1]} {origin[2]}" '
42                 f'Spacing="{dx} {dy} 1">\n')
43         f.write(f' <Piece Extent="0 {nx-1} 0 {ny-1} 0 {nz-1}">\n')
44         f.write(f' <PointData Scalars="{name}">\n')
45         f.write(f' <DataArray type="Float32" Name="{name}" format="ascii">\n')
46         for j in range(ny):
47             row = grid[j]
48             f.write(" ".join(f"{float(row[i]):.6f}" for i in range(nx)) + "\n")
49         f.write(' </DataArray>\n')
50         f.write(' </PointData>\n')
51         f.write(' <CellData/>\n')
52         f.write(' </Piece>\n')
53         f.write(' </ImageData>\n')
54     f.write('</VTKFile>\n')
55
56 def write_pvd_collection(pvd_path: str, vtk_filenames: List[str], times: List[float]):
57     """
58     Write a ParaView Data (PVD) collection file so the sequence loads as a time series.
59     """
60     assert len(vtk_filenames) == len(times)
61     with open(pvd_path, "w") as f:
62         f.write('<?xml version="1.0"?>\n')
63         f.write('<VTKFile type="Collection" version="0.1" byte_order="LittleEndian">\n')
64         f.write(" <Collection>\n")
65         for fn, t in zip(vtk_filenames, times):
66             f.write(f' <DataSet timestep="{t:.6f}" file="{Path(fn).name}">\n')
67         f.write(" </Collection>\n")
68     f.write("</VTKFile>\n")
69
70 def write_csv_snapshot(path: str, grid: "list[list[float]]"):
71     ny = len(grid)
72     nx = len(grid[0])
73     with open(path, "w") as f:
74         f.write("i,j,T\n")
75         for j in range(ny):
76             for i in range(nx):
```

```
77         f.write(f"{i},{j},{grid[j][i]:.6f}\n")  
78  
79     def create_temperature_frame(grid: "list[list[float]]", step: int, delta: float,  
80                                  vmin: float, vmax: float, dpi: int = 100,  
81                                  extent: tuple = None) -> np.ndarray:  
82         """  
83             Create a matplotlib figure of the temperature field and return it as an array.  
84             extent: (xmin, xmax, ymin, ymax) for physical dimensions  
85         """  
86         import matplotlib  
87         import matplotlib.colors as mcolors  
88         matplotlib.use('Agg') # Use non-interactive backend  
89  
90         ny = len(grid)  
91         nx = len(grid[0])  
92  
93         # Convert grid to numpy array (flip y-axis for correct orientation)  
94         T_array = np.array(grid)  
95  
96         fig, ax = plt.subplots(figsize=(8, 7), dpi=dpi)  
97  
98         # Create desaturated jet colormap to match ParaView's smooth blue-to-red gradient  
99         base_cmap = plt.cm.jet  
100        # Desaturate by blending with gray (reduces harsh transitions)  
101        def desaturate_colormap(cmap, sat=0.7):  
102            colors = cmap(np.linspace(0, 1, 256))  
103            # Convert to HSV, reduce saturation, convert back  
104            h = colors[:, 0]  
105            s = colors[:, 1] * sat # Reduce saturation  
106            v = colors[:, 2]  
107            # For simplicity, just blend with gray  
108            gray = 0.5  
109            colors[:, :3] = colors[:, :3] * sat + gray * (1 - sat)  
110            return mcolors.ListedColormap(colors)  
111  
112        desaturated_rainbow = desaturate_colormap(base_cmap, sat=0.75)  
113  
114        # Use extent for physical dimensions if provided  
115        if extent:  
116            im = ax.imshow(T_array, origin='lower', cmap=desaturated_rainbow, vmin=vmin, vmax=vmax,  
117                           aspect='equal', interpolation='bilinear', extent=extent)  
118        else:  
119            im = ax.imshow(T_array, origin='lower', cmap=desaturated_rainbow, vmin=vmin, vmax=vmax,  
120                           aspect='equal', interpolation='bilinear')  
121  
122        # Add colorbar  
123        cbar = plt.colorbar(im, ax=ax, label='Temperature (F)')  
124  
125        # Add title with iteration info  
126        ax.set_title(f'Temperature Field Evolution\nIteration: {step}, Delta: {delta:.2e}',  
127                     fontsize=12, fontweight='bold')  
128  
129        # Set axis labels based on whether physical dimensions are provided  
130        if extent:  
131            ax.set_xlabel('X Position (meters)')
```

```
132         ax.set_ylabel('Y Position (meters)')
133     else:
134         ax.set_xlabel('X Position')
135         ax.set_ylabel('Y Position')
136
137     # Convert figure to numpy array (using buffer_rgba for compatibility)
138     fig.canvas.draw()
139     buf = np.asarray(fig.canvas.buffer_rgba())
140     # Convert RGBA to RGB
141     frame = buf[:, :, :3].copy()
142
143     plt.close(fig)
144     return frame
145
146     # -----
147     # Physics / Discretization
148     # -----
149
150 def build_indices_hot_square(nx: int, ny: int, hot_fraction: float = 1/3.0) -> Tuple[int, int, int, int]:
151     """
152     Hot square is hot_fraction of plate width/height, centered.
153     We keep it integer-index aligned and inclusive bounds.
154     """
155     hot_w = int(round(nx * hot_fraction))
156     hot_h = int(round(ny * hot_fraction))
157
158     i0 = (nx - hot_w) // 2
159     i1 = i0 + hot_w - 1
160     j0 = (ny - hot_h) // 2
161     j1 = j0 + hot_h - 1
162     return i0, i1, j0, j1
163
164 def apply_dirichlet_boundaries(T: "list[list[float]]", top: float, bottom: float):
165     ny = len(T)
166     nx = len(T[0])
167
168     # Top and Bottom
169     for i in range(nx):
170         T[ny - 1][i] = top
171         T[0][i] = bottom
172
173     # Left/Right linear ramps from bottom (32) to top (100)
174     for j in range(ny):
175         y = j / (ny - 1) if ny > 1 else 0.0
176         ramp = bottom + (top - bottom) * y
177         T[j][0] = ramp
178         T[j][nx - 1] = ramp
179
180 def apply_hot_square(T: "list[list[float]]", i0: int, i1: int, j0: int, j1: int, temp: float):
181     for j in range(j0, j1 + 1):
182         for i in range(i0, i1 + 1):
183             T[j][i] = temp
184
185 def jacobi_step(T_old: "list[list[float]]", T_new: "list[list[float]]",
186                 fixed_mask: "list[list[bool]]") -> float:
```

```
187     """
188     Perform one Jacobi sweep; return delta = max |T_new - T_old|
189     fixed_mask=True means Dirichlet node (boundary or hot square): copy through.
190     """
191
192     ny = len(T_old)
193     nx = len(T_old[0])
194     delta = 0.0
195
196     for j in range(ny):
197         jm1 = max(j - 1, 0)
198         jp1 = min(j + 1, ny - 1)
199         for i in range(nx):
200             if fixed_mask[j][i]:
201                 T_new[j][i] = T_old[j][i]
202             else:
203                 im1 = max(i - 1, 0)
204                 ip1 = min(i + 1, nx - 1)
205                 T_new[j][i] = 0.25 * (T_old[j][im1] + T_old[j][ip1] +
206                                         T_old[jm1][i] + T_old[jp1][i])
207                 d = abs(T_new[j][i] - T_old[j][i])
208                 if d > delta:
209                     delta = d
210
211     # -----
212     # Main
213     # -----
214
215 def main(hot_fraction=None, return_data=False):
216     ap = argparse.ArgumentParser(description="Laplace (Jacobi) with ParaView-ready outputs.")
217     ap.add_argument("--nx", type=int, default=181, help="Grid points in x (columns.)")
218     ap.add_argument("--ny", type=int, default=181, help="Grid points in y (rows.)")
219     ap.add_argument("--tol", type=float, default=1e-3, help="Convergence tolerance on delta.")
220     ap.add_argument("--max-iters", type=int, default=20000, help="Safety cap on iterations.")
221     ap.add_argument("--output-every", type=int, default=20, help="Write VTK every N iterations.")
222     ap.add_argument("--out", type=str, default="out_vtk", help="Output directory.")
223     ap.add_argument("--also-csv", action="store_true", help="Write CSV snapshots too.")
224     ap.add_argument("--gif", action="store_true", help="Generate animated GIF of temperature evolution.")
225     ap.add_argument("--gif-fps", type=int, default=10, help="Frames per second for GIF.")
226     ap.add_argument("--hot-fraction", type=float, default=1/3.0, help="Hot square size as fraction of domain
227     (default: 1/3)")
228     ap.add_argument("--no-vtk", action="store_true", help="Skip VTK file writing (for webapp use)")
229     args = ap.parse_args()
230
231     # Override with programmatic argument if provided
232     if hot_fraction is not None:
233         args.hot_fraction = hot_fraction
234
235     # Physical sizes (meters)
236     W = 9.0
237     H = 9.0
238     dx = W / (args.nx - 1)
239     dy = H / (args.ny - 1)
240
241     # Temperatures (degrees)
```

```
241     T_BOTTOM = 32.0
242     T_TOP = 100.0
243     T_HOT = 212.0
244
245     out_dir = Path(args.out)
246     out_dir.mkdir(parents=True, exist_ok=True)
247
248     # Allocate fields
249     T_old = [[T_BOTTOM for _ in range(args.nx)] for _ in range(args.ny)]
250     T_new = [[T_BOTTOM for _ in range(args.nx)] for _ in range(args.ny)]
251
252     # Apply boundaries and hot square to initial field
253     apply_dirichlet_boundaries(T_old, T_TOP, T_BOTTOM)
254
255     i0, i1, j0, j1 = build_indices_hot_square(args.nx, args.ny, args.hot_fraction)
256     apply_hot_square(T_old, i0, i1, j0, j1, T_HOT)
257
258     # Build fixed mask (Dirichlet everywhere that must not change)
259     fixed = [[False for _ in range(args.nx)] for _ in range(args.ny)]
260     # Outer boundaries:
261     for i in range(args.nx):
262         fixed[0][i] = True
263         fixed[args.ny - 1][i] = True
264     for j in range(args.ny):
265         fixed[j][0] = True
266         fixed[j][args.nx - 1] = True
267     # Hot square region:
268     for j in range(j0, j1 + 1):
269         for i in range(i0, i1 + 1):
270             fixed[j][i] = True
271
272     # Save step 0
273     vtk_paths = []
274     times = []
275     frames = [] if args.gif else None
276     step = 0
277
278     if not args.no_vtk:
279         vti0 = out_dir / f"step_{step:05d}.vti"
280         write_vti_image_data(str(vti0), T_old, dx, dy, name="Temperature")
281         vtk_paths.append(str(vti0))
282         times.append(float(step))
283         if args.also_csv:
284             write_csv_snapshot(str(out_dir / f"step_{step:05d}.csv"), T_old)
285     if args.gif:
286         # Create extent for 9m x 9m physical dimensions
287         extent = (0, W, 0, H)
288         frames.append(create_temperature_frame(T_old, step, 0.0, T_BOTTOM, T_HOT, extent=extent))
289
290     # Jacobi iteration loop (slide: do { ... } while (delta > tolerance))
291     delta = float("inf")
292     convergence_history = {"iterations": [], "deltas": []}
293
294     while delta > args.tol and step < args.max_iters:
295         delta = jacobi_step(T_old, T_new, fixed)
```

```
296     T_old, T_new = T_new, T_old # swap buffers
297     step += 1
298
299     # Track convergence history
300     convergence_history["iterations"].append(step)
301     convergence_history["deltas"].append(delta)
302
303     # Output cadence like the slides' "t=0,5,10,20,..."
304     if step % args.output_every == 0 or delta <= args.tol:
305         if not args.no_vtk:
306             vti_path = out_dir / f"step_{step:05d}.vti"
307             write_vti_image_data(str(vti_path), T_old, dx, dy, name="Temperature")
308             vtk_paths.append(str(vti_path)) # keep the list; it now holds .vti files
309             times.append(float(step))
310             if args.also_csv:
311                 write_csv_snapshot(str(out_dir / f"step_{step:05d}.csv"), T_old)
312                 print(f"[iter {step:6d}] delta={delta:.6e} -> wrote {vti_path.name}")
313         else:
314             print(f"[iter {step:6d}] delta={delta:.6e}")
315     if args.gif:
316         extent = (0, W, 0, H)
317         frames.append(create_temperature_frame(T_old, step, delta, T_BOTTOM, T_HOT, extent=extent))
318
319     # PVD collection for ParaView time series
320     if not args.no_vtk and len(vtk_paths) > 0:
321         write_pvd_collection(str(out_dir / "series.pvd"), vtk_paths, times)
322
323     # Save GIF if requested
324     if args.gif and len(frames) > 0:
325         gif_path = out_dir / "temperature_evolution.gif"
326         print(f"\nSaving GIF animation with {len(frames)} frames...")
327
328         # Use PIL to save GIF
329         from PIL import Image
330         pil_frames = [Image.fromarray(frame) for frame in frames]
331         duration_ms = int(1000 / args.gif_fps)
332         pil_frames[0].save(
333             str(gif_path),
334             save_all=True,
335             append_images=pil_frames[1:],
336             duration=duration_ms,
337             loop=0
338         )
339         print(f"GIF saved to: {gif_path}")
340
341     print(f"\nDone. Final iter={step}, delta={delta:.6e}")
342     if not args.no_vtk:
343         print(f"Open in ParaView: {out_dir/'series.pvd'} (time series)")
344         print("Tip: add 'Annotate Time' for the iteration number, and show the scalar bar.")
345     if args.gif:
346         print(f"View animation: {out_dir/'temperature_evolution.gif'}")
347
348     # Return data if called programmatically
349     if return_data:
350         return {
```

```
351     "convergence_history": convergence_history,
352     "final_grid": T_old,
353     "final_iter": step,
354     "final_delta": delta,
355     "hot_fraction": args.hot_fraction,
356     "nx": args.nx,
357     "ny": args.ny,
358     "frames": frames if args.gif else None
359   }
360
361 if __name__ == "__main__":
362   main()
```